

# Online Geometric Optimization in the Bandit Setting Against an Adaptive Adversary

H. Brendan McMahan and Avrim Blum

Carnegie Mellon University


December 21, 2024

# Introduction

- Overview of online optimization in adversarial environments.
- Importance of studying bandit settings with limited feedback.
- Addressing challenges posed by adaptive adversaries.


# Context: Online Optimization with an Adversary

- Online optimization involves making sequential decisions.
- Adversaries can influence the environment by selecting cost vectors.
- Applications in machine learning, economics, and network routing.



# Objective: Solving the Bandit Version

- Transition from full-information to bandit feedback.
- Challenges due to limited information after each decision.
- Goal: Develop algorithms that perform well despite these limitations.



# Description of the Online Optimization Problem

- Set of feasible points  $S \subset \mathbb{R}^n$ .
- Sequential decision-making over  $T$  rounds.
- At each round  $t$ , select  $x_t \in S$  and incur cost  $c_t \cdot x_t$ .

## Mathematical Example


**Example:** Let  $S$  be the unit simplex in  $\mathbb{R}^n$ , i.e.,

$$S = \left\{ x \in \mathbb{R}^n \mid x_i \geq 0 \text{ for all } i, \sum_{i=1}^n x_i = 1 \right\}$$

This setup is common in online portfolio selection.

# Adversary and Algorithm Interactions

- Adversary selects cost vectors  $c_t \in \mathbb{R}^n$ .
- Algorithm selects decisions  $x_t \in S$  without knowing  $c_t$  in advance.
- Objective is to minimize cumulative cost over  $T$  rounds.



- $S$  represents all possible feasible decisions.
- Geometric properties of  $S$  influence algorithm design.
- Examples: Convex sets, polytopes, or combinatorial structures.

## Mathematical Example

**Convex Set:** Let  $S = \{x \in \mathbb{R}^n \mid \|x\|_2 \leq 1\}$  be the unit ball in  $\mathbb{R}^n$ . This convex set allows the use of gradient-based optimization methods due to its smooth boundary.

## Cost Incurred: $c_t \cdot x_t$

- Inner product representing the cost for decision  $x_t$ .
- $c_t$  is the cost vector chosen by the adversary.
- Objective: Minimize the sum  $\sum_{t=1}^T c_t \cdot x_t$ .

$$\text{Total Cost} = \sum_{t=1}^T c_t \cdot x_t$$



# The Role of an Oracle for Offline Optimization

- Oracle solves  $\min_{x \in S} (c \cdot x)$  efficiently.
- Assumes full knowledge of cost vectors.
- Serves as a benchmark for online algorithms.

## Mathematical Example

**Offline Optimization:** Given all cost vectors  $\{c_1, c_2, \dots, c_T\}$  in advance, the oracle solves:

$$\min_{x \in S} \sum_{t=1}^T c_t \cdot x$$

This provides the best possible cumulative cost.

Objective:  $\min_{x \in S} (c \cdot x)$

- Offline problem assumes all cost vectors are known in advance.
- Provides the optimal benchmark for comparison.
- Online algorithms aim to perform nearly as well without this knowledge.

$$\text{Optimal Offline Cost} = \min_{x \in S} \sum_{t=1}^T c_t \cdot x$$

# Observing Only the Cost: $c_t \cdot x_t$

- Limited feedback: Only the incurred cost is observed.
- No access to the full cost vector  $c_t$ .
- Necessitates exploration to gather information.

## Concrete Example

**Example:** In online advertising, selecting an ad (decision  $x_t$ ) and only observing the click-through rate (incurred cost) without knowing the underlying user preferences (cost vector  $c_t$ ).

# Example: Online Shortest Path Problem

- Decision set  $S$  consists of all possible paths in a network.
- Cost vectors represent edge weights chosen by the adversary.
- Feedback is the total cost of the chosen path only.

## Concrete Example

Consider a network with 4 nodes (A, B, C, D) and 5 edges. At each round, an adversary assigns weights to the edges. The algorithm selects a path from node A to node D and only observes the sum of the weights on the chosen path.

# Difference Between Oblivious and Adaptive Adversaries

- **Oblivious Adversary:** Chooses all cost vectors in advance.
- **Adaptive Adversary:** Chooses  $c_t$  based on past algorithm decisions.
- Adaptive adversaries are more powerful and challenging.

Adversary Type	Decision
Oblivious	Chooses $c_t$ in advance
Adaptive	Chooses $c_t$ based on past decisions

**Table:** Comparison of Adversary Types

# Impact of Adversary's Strategy on the Algorithm

- Adaptive strategies can exploit algorithm's weaknesses.
- Necessitates robust algorithms that can handle changing environments.
- Importance of maintaining low regret despite adversary's adaptability.

# Formal Definition of Regret

- **Regret** measures the performance gap between the algorithm and the best offline decision.
- Mathematically:

$$\text{Regret} = \mathbb{E} \left[ \sum_{t=1}^T c_t \cdot x_t \right] - \mathbb{E} \left[ \min_{x \in S} \sum_{t=1}^T c_t \cdot x \right]$$

- $\mathbb{E}$ : Expectation over the algorithm's randomness.
- $c_t$ : Cost vector at round  $t$ .
- $x_t$ : Decision made by the algorithm at round  $t$ .
- Goal is to achieve sublinear regret:  $\text{Regret} = o(T)$ .

# Goal: Minimize Regret

- Achieve sublinear regret:  
 $\text{Regret} = o(T)$ .
- Ensures average regret per round goes to zero as  $T$  increases.
- Fundamental objective in online learning and optimization.



# Description of the Kalai-Vempala Algorithm

- General algorithm for online convex optimization.
- Utilizes the exponential weights framework.
- Assumes full feedback: Observes entire cost vector  $c_t$  after each decision.

## Mathematical Formulation

The Kalai-Vempala algorithm updates the weight  $w_t(x)$  for each decision  $x \in S$  as:

$$w_{t+1}(x) = w_t(x) \exp(-\eta c_t \cdot x)$$

where  $\eta$  is the learning rate.

# Assumption: Adversary Selects Cost Vectors After Decision

- The algorithm selects  $x_t$  without knowing  $c_t$ .
- After selection,  $c_t$  is revealed.
- Enables the use of gradient-based updates in the algorithm.

$$x_t = \frac{w_t(x)}{\sum_{x' \in S} w_t(x')}$$

# Transition to the Bandit Setting

- In the bandit setting, only the incurred cost  $c_t \cdot x_t$  is observed.
- No access to the full cost vector  $c_t$ .
- Requires estimation techniques to infer  $c_t$  from limited feedback.

## Estimation Technique

**Importance Sampling:** To estimate the cost vector, the algorithm can use techniques like importance sampling where:

$$\hat{c}_t = \frac{c_t \cdot x_t}{P(x_t)} e_{x_t}$$

where  $P(x_t)$  is the probability of selecting  $x_t$  and  $e_{x_t}$  is the standard basis vector corresponding to  $x_t$ .

# Key Challenge: Working with Partial Feedback


- Limited information makes it difficult to update cost estimates accurately.
- Balancing exploration (gathering information) and exploitation (using current knowledge).
- Ensuring robust performance against adaptive adversaries.
- Incorporating additional mathematical tools for estimation.

# Description of the Bandit-style Geometric Decision Algorithm (BGA)

- Designed for bandit feedback in online geometric optimization.
- Combines exploration and exploitation phases.
- Leverages geometric properties of the decision set  $S$ .

# Alternating Between Exploration and Exploitation

- **Exploitation:** Use current estimates to make informed decisions.
- **Exploration:** Sample randomly from a basis to gather new information.
- Ensures that the algorithm learns the cost structure over time.




# Pseudocode for BGA Algorithm

## BGA Algorithm

- 1 Initialize estimates and sampling basis  $B \subset S$ .
- 2 For each round  $t = 1$  to  $T$ :
  - 1 With probability  $\gamma$ , **explore** by selecting a random basis vector  $b \in B$ .
  - 2 Otherwise, **exploit** by selecting  $x_t$  using the Kalai-Vempala strategy on estimated costs.
  - 3 Observe the incurred cost  $c_t \cdot x_t$ .
  - 4 Update cost estimates based on the observed cost.

## Flowchart Representation



# Explanation of the Sampling Basis $B \subset S$

- $B$  is a carefully chosen subset of  $S$  that facilitates exploration.
- Ensures coverage of different regions in the decision space.
- Basis vectors are used to approximate the cost structure.
- Geometric properties of  $B$  influence the efficiency of exploration.

## Mathematical Insight

If  $B$  forms a **basis** in the linear algebra sense, any decision  $x \in S$  can be expressed as a linear combination of vectors in  $B$ . This allows efficient reconstruction and estimation of the cost vector  $c_t$ .



# Decision Making Based on Basis $B$

- During exploration, a basis vector  $b$  is selected uniformly at random.
- Provides unbiased estimates of the cost vector  $c_t$ .
- Facilitates efficient updating of cost estimates.

$$\hat{c}_t = \frac{c_t \cdot b}{P(b)} e_b$$

where  $P(b)$  is the probability of selecting basis vector  $b$ , and  $e_b$  is the corresponding basis vector.

# Exploration Probability $\gamma$

- $\gamma$  controls the frequency of exploration.
- Higher  $\gamma$  leads to more exploration, improving cost estimates.
- Lower  $\gamma$  emphasizes exploitation, leveraging current knowledge.

## Balancing $\gamma$

The optimal choice of  $\gamma$  depends on problem parameters such as  $T$  and  $n$ . Typically,  $\gamma$  is set to decrease over time to balance exploration and exploitation effectively.

# Trade-off Between Exploration and Exploitation

- Balancing  $\gamma$  is crucial for minimizing regret.
- Too much exploration can waste resources, while too little can lead to poor estimates.
- Optimal  $\gamma$  depends on problem parameters like  $T$  and  $n$ .

# Update Rule for Cost Vector Estimates

- Use observed costs to refine estimates of  $c_t$ .
- Employ matrix inversion techniques to handle correlated estimates.
- Ensures accurate approximation of the true cost vectors over time.
- Utilizes geometric properties of the decision set for efficient updates.

$$\hat{c}_{t+1} = \hat{c}_t - \eta \nabla L_t(x_t)$$

where  $\eta$  is the learning rate and  $L_t(x_t)$  is the loss at round  $t$ .

# Introduction to the Mathematical Analysis of BGA

- Formalize the algorithm's update rules and decision-making process.
- Define assumptions and properties of the cost vectors and decision set.
- Set the stage for deriving regret bounds.

## Assumptions

- 1 The cost vectors  $c_t$  are bounded, i.e.,  $\|c_t\| \leq C$  for some constant  $C$ .
- 2 The decision set  $S$  is convex and compact.
- 3 The adversary is adaptive, choosing  $c_t$  based on past decisions  $x_1, \dots, x_{t-1}$ .

# Regret Bounds and Performance Guarantees

- Analyze the cumulative regret over  $T$  rounds.
- Provide theoretical guarantees under adaptive adversary models.
- Compare performance with existing algorithms like Kalai-Vempala.

## Theoretical Insight

Under the adaptive adversary model, BGA maintains a regret bound of:

$$\text{Regret} = O\left(T^{3/4}\sqrt{\ln T}\right)$$

This ensures that the algorithm performs competitively even in dynamic environments.

# Theoretical Regret Result

- BGA achieves regret  $O\left(T^{3/4}\sqrt{\ln T}\right)$ .
- Sublinear growth ensures that average regret per round diminishes.
- Demonstrates effectiveness in the bandit setting against adaptive adversaries.

## Mathematical Proof Sketch

The regret bound is derived using:

- 1 Concentration inequalities to bound estimation errors.
- 2 Matrix inversion to handle dependencies in cost estimates.
- 3 Balancing exploration and exploitation through parameter tuning.

# Impact of Parameters $\gamma$ , $\epsilon$ , and $T$

- $\gamma$ : Balances exploration and exploitation.
- $\epsilon$ : Determines precision of cost estimates.
- $T$ : Number of rounds influences overall regret.
- Optimal tuning of parameters is essential for best performance.

$$\gamma = T^{-1/4}, \quad \epsilon = T^{-1/2}$$



# High-Probability Bounds on Cost Vector Estimates

- Establish bounds that hold with high probability.
- Use concentration inequalities to ensure reliable estimates.
- Critical for guaranteeing low regret in adversarial settings.

$$\Pr(\|\hat{c}_t - c_t\| \geq \delta) \leq \exp(-k\delta^2)$$

where  $k$  is a constant depending on  $T$  and  $n$ .

# Using Martingale Inequalities for Estimating $c_t$

- Apply martingale-based techniques to handle dependencies over time.
- Ensure that estimates remain unbiased and concentrated around true values.
- Facilitates robust analysis against adaptive adversaries.

## Theorem (Azuma-Hoeffding Inequality)

Let  $\{X_t\}$  be a martingale with bounded differences  $|X_t - X_{t-1}| \leq c_t$ .  
Then, for any  $\lambda > 0$ ,

$$\Pr(X_T - X_0 \geq \lambda) \leq \exp\left(-\frac{\lambda^2}{2 \sum_{t=1}^T c_t^2}\right)$$

# Random Exploration for Estimating True Cost Vectors

- Randomly selecting basis vectors aids in uncovering the cost structure.
- Ensures diverse coverage of the decision space.
- Reduces bias in cost estimates by providing varied perspectives.

$$\mathbb{E}[\hat{c}_t] = c_t$$

## Unbiased Estimation

The estimator  $\hat{c}_t$  is unbiased because:

$$\mathbb{E}[\hat{c}_t] = \sum_{b \in B} P(b) \cdot \frac{c_t \cdot b}{P(b)} e_b = c_t$$

# Importance of Unbiased Estimates

- Unbiased estimates are crucial for accurate decision-making.
- Prevents systematic errors that could be exploited by the adversary.
- Enhances the reliability and performance of the BGA algorithm.

$$\mathbb{E}[\hat{c}_t] = c_t$$

## Consequence

Ensures that the algorithm's decisions are based on accurate representations of the cost vectors, leading to effective optimization over time.

# Detailed Analysis of Expected Regret for BGA

- Derive bounds on the expected cumulative regret.
- Show how BGA maintains low regret despite adaptive adversaries.
- Compare theoretical performance with empirical observations.

## Expected Regret Bound

The expected regret of BGA satisfies:

$$\mathbb{E}[\text{Regret}] \leq O\left(T^{3/4}\sqrt{\ln T}\right)$$

- The bound holds under the adaptive adversary model.
- Demonstrates sublinear growth, ensuring diminishing average regret.

# Summary of Performance Guarantees

- BGA achieves sublinear regret in the bandit setting.
- Provides robustness against adaptive adversaries.
- Maintains competitive performance compared to full-information algorithms.

## Key Takeaways

- Sublinear regret ensures that the algorithm becomes more effective over time.
- Robustness against adaptive adversaries makes BGA suitable for dynamic environments.
- The bandit setting presents unique challenges that BGA effectively addresses.

# Comparison with Kalai-Vempala Algorithm

- Kalai-Vempala assumes full feedback, achieving different regret bounds.
- BGA extends these ideas to the bandit setting with limited feedback.
- Demonstrates comparable performance with additional challenges.

Algorithm	Feedback	Regret Bound
Kalai-Vempala	Full Information	$O(\sqrt{T})$
BGA	Bandit	$O(T^{3/4}\sqrt{\ln T})$

Table: Comparison of Regret Bounds

# Recap of the Main Contribution

- Introduction of BGA for online geometric optimization in bandit settings.
- Achieves low-regret performance against adaptive adversaries.
- Bridges the gap between full-information and bandit feedback models.

## Implications

BGA provides a robust framework for decision-making in environments where feedback is limited and adversaries can adapt, making it applicable to various real-world scenarios.



# Theoretical Bounds for Performance Against an Adaptive Adversary

- Established  $O(T^{3/4}\sqrt{\ln T})$  regret bound.
- Guarantees hold under general adversarial conditions.
- Highlights the algorithm's robustness and efficiency.

## Future Directions

Further research can explore tightening these bounds and extending BGA to other optimization settings.

# Discussion of Open Questions and Potential Future Work

- Can regret bounds be further improved for adaptive adversaries?
- Exploration of different sampling strategies for basis selection.
- Extensions to other types of online optimization problems.
- Incorporating additional feedback mechanisms to enhance performance.

# Improving Regret Bounds for Adaptive Adversaries

- Investigate tighter analyses and alternative algorithms.
- Explore the impact of different geometric properties of  $S$ .
- Consider hybrid models combining full and bandit feedback.
- Study the interplay between exploration rates and adversary adaptability.

## Potential Approaches

- 1 Utilize advanced concentration inequalities.
- 2 Incorporate adaptive learning rates.
- 3 Leverage multi-armed bandit techniques.

# Key References

- Kalai, E., & Vempala, S. (2005). A polynomial time algorithm for online convex programming. *Proceedings of the ACM Conference on Learning Theory*, 161-170.
- McMahan, H. B., & Blum, A. (Year). *Title of the Paper*. *Journal/Conference Name*.
- Auer, P., Cesa-Bianchi, N., & Fischer, P. (2002). Finite-time Analysis of the Multiarmed Bandit Problem. *Machine Learning*, 47(2-3), 235-256.

ithms.

# Acknowledgments

- Thank collaborators and contributors.
- Acknowledge funding sources and institutional support.
- Express gratitude to reviewers and advisors.

## Example

We would like to thank our colleagues at Carnegie Mellon University for their invaluable feedback and support. This work was supported by NSF Grant #XXXXXX.